

Computer Architecture. Week 9

Muhammad Fahim, Alexander Tormasov

Innopolis University

m.fahim@innopolis.ru

a.tormasov@innopolis.ru

October 25, 2020

Topic of the lecture

- The Processor

Topic of the tutorial

- Mid-term Exam

Topic of the lab

- The Single Cycle Processor

Content of the class

- Datapath and Control Signals
- Architecture Implementations
- The Building Blocks of Processor
- Steps in Designing the Processor
- Memories
- Pipelining
- Stages of Pipeline
- Pipeline Datapath and Control Signal

Datapath and Control

- How to implement an architecture in hardware?
- **Processor**
 - **Datapath:** A datapath is a collection of functional units – such as arithmetic logic unit that performs data processing operations – registers, and buses.
 - **Control:** Control signals

Architecture Implementations

- Multiple implementations for a single architecture
 - **Single-cycle:** All operations take the same amount of time – a single cycle
 - **Multicycle:** It allows faster operations to take less time than slower ones, so overall performance can be increased
 - **Pipelined:** Each instruction is broken up into series of steps multiple instructions execute at once

Recap: Processor Performance


- Program execution time

$$\text{CPU Time} = \# \text{instructions} \times \text{CPI} \times \text{Clock cycle time}$$

- CPU performance factors
 - Instruction count – Determined by Instruction Set Architecture (ISA) and compiler
 - CPI and cycle time – Determined by implementation of the processor
- Challenge is to satisfy constraints of
 - Cost
 - Power
 - Performance

Computers are state machines

- Computer is just a big fancy state machine.
- The processor keeps reading and updating the state, according to the instructions in some program.



Architectural State

- Determines everything about a processor
 - PC (Program Counter)
 - 32 Registers
 - Memory


Note: Architectural state includes register files and memory, and microarchitectural state represents **internal state that has not yet been exposed outside the processor** (For example: instruction queue state).

MIPS Processor

- Consider subset of MIPS instructions
 - Memory-reference instructions: `lw`, `sw`
 - R-type instructions: `add`, `sub`, `and`, `or`, `slt`
 - Branch instructions: `beq`, `bne`

Datapath

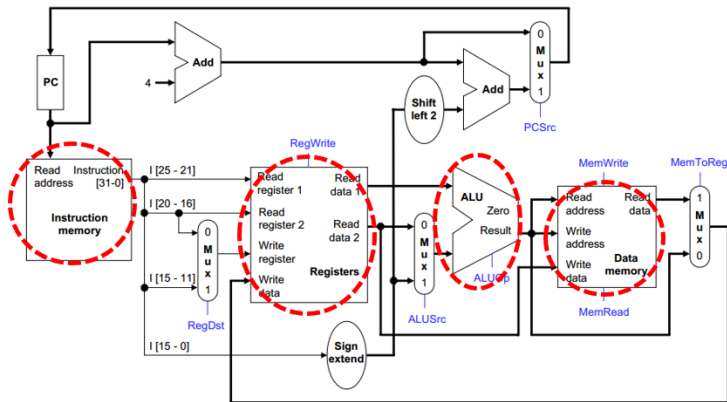
- Simplified view:



NOTE:

PC is Program Counter. It is also known as instruction pointer.


- Final view:



- NOTE: More details in tutorial session.

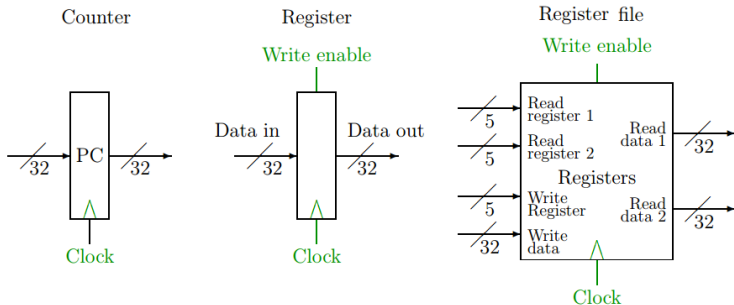
The Building Blocks of Processor (1/2)

- We have already designed many of the major components for the processor, or have at least identified how they could be implemented.



- Note: The diagram highlights the control signals in green color.

The Building Blocks of Processor (2/2)



Steps in Designing a Processor

- Register Transfer Language (RTL) is a kind of intermediate representation that is very close to assembly language.
- It is used to describe data flow at the **register-transfer level** of an architecture.
- From the RTL description of each instruction, determine
 - The required datapath components
 - The datapath interconnections
- Determine the control signals required to enable the datapath elements in the appropriate sequence for each instruction.
- **Usually, it is not exposed to programmers.**

Register Transfer Level (RTL)

- It is a design abstraction and used to describe instructions.
- It models the flow of data between registers.
- Any instruction cycle starts by fetching the instruction.
- Typically, every instruction involves operations and modification of the PC.

RTL Example

- The ADD instruction: `add rd, rs, rt`

 1. `mem[PC]` Fetch the instruction from memory
 2. `R[rd] ← R[rs] + R[rt]` Set register `rd` to the value of the sum of the contents of registers `rs` and `rt`
 3. `PC ← PC + 4` Calculate the address of the next instruction
- NOTE: More details in tutorial session.

Remarks! Datapath and Control

- A datapath contains **all the functional units** and **necessary connections** to implement an instruction set architecture.
- The **control unit** tells the datapath what to do, based on the instruction that's currently being executed.

John von Neumann (1/2)

- In the old days, “programming” involved actually **changing a machine’s physical configuration by flipping switches or connecting wires.**
 - A computer could run just one program at a time.
 - Memory only stored data that was being operated on.


John von Neumann (2/2)

- Then around 1944, John von Neumann and others got the idea to encode instructions in a format that could be stored in memory just like data.
 - The processor interprets and executes instructions from memory.
 - One machine could perform many different tasks, just by loading different programs into memory.
 - The “stored program” design is often called a Von Neumann machine

- In modern computers, programs are stored in the “text segment” of memory, which can be written to only by the operating system.

Memories: Harvard Architecture

- It's easier to use a **Harvard architecture** at first, with programs and data stored in **separate memories**.




Multicycle Processor

- A multicycle processor **fixes some shortcomings** in the single-cycle CPU.
 - Faster instructions are not held back by slower ones.
 - The clock cycle time can be decreased.
 - A multicycle processor requires a somewhat simpler datapath


Pipelining (1/4)

Laundry Example

- Ann, Brian, Cathy, Dave each have one load of clothes to wash, dry, and fold.
- Washer takes 30 minutes
- Dryer takes 40 minutes
- “Folder” takes 20 minutes




Pipelining (2/4)



- Sequential laundry takes 6 hours for 4 loads
- If they learned pipelining, how long would laundry take?


Pipelining (3/4)



- Pipelined laundry takes 3.5 hours for 4 loads

Pipelining (4/4)

- Pipelining doesn't help latency of single task, it helps throughput of entire workload



Pipeline Concept in Computer Architecture

- Break up instruction into tasks
- Balance the amount of work (time) between stages
- Allow each segment to complete and start next instruction

Properties of Pipelines


- Latency:** The time it takes for a single instruction to execute. Pipelining makes latency slightly worse.
- Throughput:** The number of instructions executed per unit time. Pipelining improves throughput.
- Pipelining Idea:** The idea behind pipelining is to maximize the usage of the available resources by overlapping the execution of several tasks.

Pipelining a Digital System (1/2)

- Key idea: Break big computation up into pieces



- Separate each piece with a pipeline register



Pipelining a Digital System (2/2)

- The idea behind pipelining is to maximize the usage of the hardware by overlapping the execution of several instructions.

Stages of Pipeline

- Five stages in classical pipeline
 - Stage 1: Instruction Fetch (IF)
 - Stage 2: Instruction Decode (ID)
 - Stage 3: Execute (EX)
 - Stage 4: Memory (MEM)
 - Stage 5: Write Back (WB)
- Clock is constrained by slowest stage of pipeline.
- Pipelining is not free: complexities in design and additional resources (more later).

Stage 1: Instruction Fetch (IF)

- Fetch an instruction from memory at every cycle
 - Use **Program Counter** (PC) to index memory
 - Increment PC (assume no branches for now)
- Write state to the pipeline register (IF/ID)
 - The next stage will read this pipeline register

Stage 2: Instruction Decode (ID)

- ◉ Decodes operation code (opcode) bits
 - ◉ Set up Control signals for later stages
- ◉ Read input operands from register file
 - ◉ Specified by decoded instruction bits
- ◉ Write state to the pipeline register (ID/EX)
 - ◉ Opcode
 - ◉ Register contents
 - ◉ $PC + 4$
 - ◉ Control signals

Stage 3: Execute (EX)

- Perform ALU operations
 - Calculate result of instruction
 - Control signals select operation
 - Contents of register 1 used as one input
 - Either register 2 or constant offset (from instruction) used as second input
- Calculate PC-relative branch target (if any)
 - $PC + 4 + (\text{constant offset})$
- Write state to the pipeline register (EX/Mem)
 - ALU result, contents of register 2, and $PC + 4 + \text{offset}$
 - Control signals

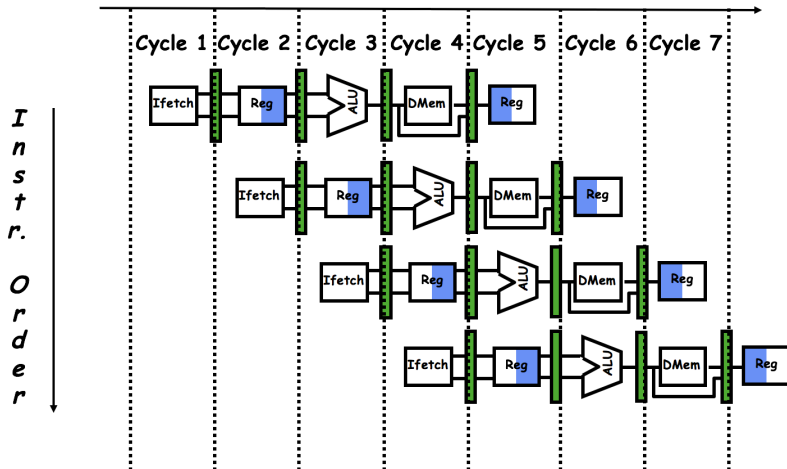
Stage 4: Memory (MEM)

- Perform data operations
 - ALU result contains address for `ld` or `sw`
 - Operation code bits control read/write and enable signals
- Write state to the pipeline register (Mem/WB)
 - ALU result and loaded data
 - Control signals


Stage 5: Writeback (WB)

- Writing result to register file (if required)
 - Write loaded data to destination register for `lw`
 - Write ALU result to destination register for arithmetic instruction
 - Operation code bits control register write enable signal

The Basic Pipeline For MIPS



A Pipeline Diagram



- A pipeline diagram shows the execution of a series of instructions.

 - The instruction sequence is shown vertically, from top to bottom.
 - Clock cycles are shown horizontally, from left to right.
 - Each instruction is divided into its component stages.
- This clearly indicates the overlapping of instructions. For example, there are three instructions active in the third cycle above.

 - lw instruction is in its execute stage.
 - sub is in its Instruction decode stage.
 - and instruction is just being fetched.

Pipeline Terminology

	1	2	3	4	5	6	7	8	9
lw \$t0, 4(\$sp)	IF	ID	EX	MEM	WB				
sub \$v0, \$a0, \$a1		IF	ID	EX	MEM	WB			
and \$t1, \$t2, \$t3			IF	ID	EX	MEM	WB		
or \$s0, \$s1, \$s2				IF	ID	EX	MEM	WB	
addi \$sp, \$sp, -4					IF	ID	EX	MEM	WB


- The **pipeline depth** is the number of stages – in this case, five.
- In the first four cycles here, the **pipeline is filling**, since there are unused functional units.
- In **cycle 5**, the **pipeline is full**. Five instructions are being executed simultaneously, so all hardware units are in use.
- In cycles 6-9, the **pipeline is emptying**

Exceptions in Pipelining


- Pipeline provides ability for processor to handle exception, save state, and restart **without affecting program execution**.
- Pipeline is **restartable**
- All processors support this feature now, as **it is needed to implement virtual memory**

Single-Cycle vs. Pipeline Execution

Non-Pipeline



Pipeline



What about Control Signals?

- The **control signals** are generated in the same way as in the **single-cycle processor** – after an instruction is fetched, the processor decodes it and produces the appropriate control values.
- But just like before, **some of the control signals will not be needed until some later stage and clock cycle.**
- These signals must be **propagated through the pipeline until they reach the appropriate stage.** We can just pass them in the pipeline registers, along with the other data.

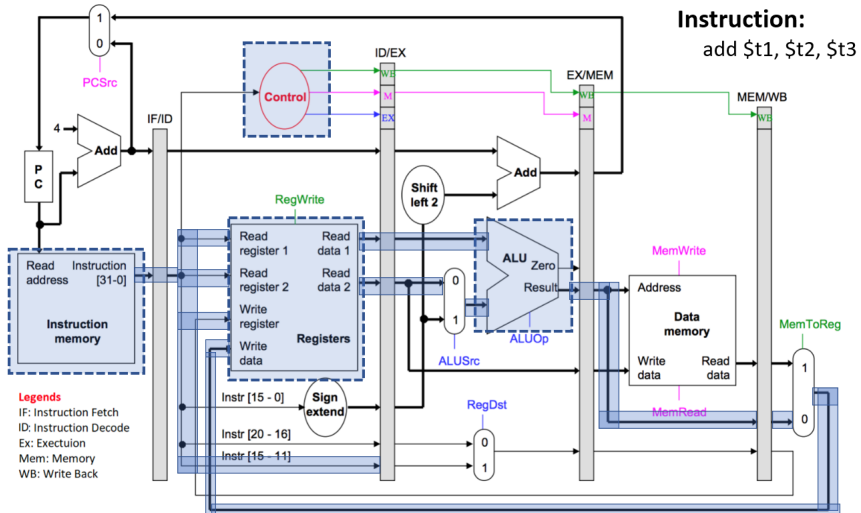
What about Control Signals?

- Control signals can be categorized by the pipeline stage

Stage	Control signal needed		
EX	ALUSrc	ALUOp	RegDst
MEM	MemRead	MemWrite	PCSrc
WB	RegWrite	MemToReg	

Note: The details will be discussed in tutorial.

Pipeline Datapath and Control Signal



Remarks on Pipelining (1/2)

- Pipelining attempts to maximize hardware usage by overlapping the execution stages of several different instructions.
- Pipelining offers amazing speedup.
 - The CPU throughput is dramatically improved, because several instructions can be executing concurrently.
 - In the best case, one instruction finishes on every cycle, and the speedup is equal to the pipeline depth.

Remarks on Pipelining (2/2)

- The bad news
 - Instructions can interfere with each other **hazards**
 - It may increase the latency
 - Different instructions may need the same piece of hardware (e.g., memory) in same clock cycle
 - **For Example:** Instruction may require a result produced by an earlier instruction that is not yet complete
- All these details are in next lecture!!

Summary

- Datapath and Control Signals
- Architecture Implementations
- The Building Blocks of Processor
- Steps in Designing the Processor
- Memories
- Pipelining and its stages
- Pipeline Datapath and Control Signal

Acknowledgements

- This lecture was created and maintained by Muhammad Fahim, Giancarlo Succi and Alexander Tormasov