

Computer Architecture. Week 10

Muhammad Fahim, Alexander Tormasov

Innopolis University

m.fahim@innopolis.ru

a.tormasov@innopolis.ru

November 05, 2020

Topic of the lecture

- Hazards and its solutions

Topic of the tutorial

- Simple ALU Unit

Topic of the lab

- Adder implementations

Content of the class

- Recap: Pipelining
- Hazards
- Types of Hazards
- Structural Hazards and its Solution
- Data Hazards and its Solution
- Control Hazards and its Solution
- Summary

Recap: Pipelining (1/2)

- Pipelining attempts to **maximize hardware usage** by overlapping the execution stages of several different instructions.
- Pipelining offers **amazing speedup**.
 - The CPU throughput is dramatically improved, because several instructions can be executing concurrently.
 - In the best case, one instruction finishes on every cycle, and the speedup is equal to the pipeline depth.


Recap: Pipelining (2/2)

- The bad news: Instructions can interfere with each other – known as **hazards**
- Different instructions may need the same piece of hardware (e.g., memory) in same clock cycle
- For Example: Instruction may require a result produced by an earlier instruction that is not yet complete

Hazards

- Hazards prevent next instruction from execution during its designated clock cycle
- Hazards **reduce the performance** from the ideal speedup gained by pipeline
- Three types of hazards
 - Structural hazards
 - Data hazards
 - Control hazards

Types of Hazards and Solutions



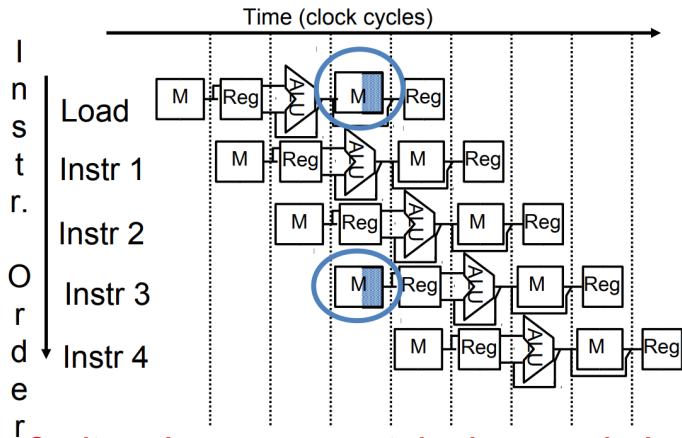
- GPR: General Purpose Registers

Structural Hazards (1/3)

- **Structural Hazards:** Attempt to use the same resource by two or more instructions at the same time
- **Example:** Single Memory for instructions and data
 - Accessed by Instruction Fetch (IF) stage
 - Accessed at same time by Memory (MEM) stage
- **Solutions**
 - Delay the second access by one clock cycle
 - **OR**
 - Provide separate memories for instructions and data
 - This is called a “Harvard Architecture”
 - Real pipeline processors have separate **caches**

Structural Hazards (2/3)

- Consider a Von Neumann architecture (same memory for instructions and data)



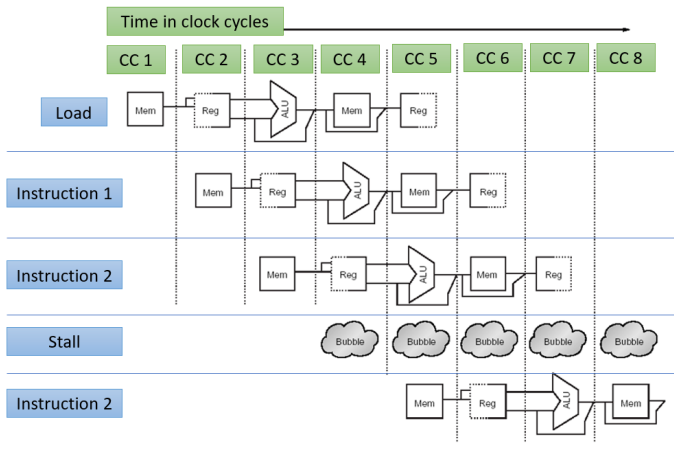
Can't read same memory twice in same clock cycle

Structural Hazards (3/3)

● Explanation

- The load instruction wants to access the memory to load data.
- At the same time instruction 3 wants to fetch an instruction from memory.

Structural Hazards – Solution (1/2)



- Stall cycle added (commonly called pipeline bubble)
- A machine **without** structural hazard will have lower CPI

Structural Hazards – Solution (2/2)


- Why a designer allows structural hazard?
- **To reduce cost**
 - Pipelining all the functional units or duplicating them may be too costly
- **To reduce latency**
 - Making functional units pipelined adds delay (pipeline overhead \rightarrow registers.)

Data Hazards (1/4)

- **Data Hazards:** It attempts to use data before it is ready
- Consider the execution of following instructions, on our pipeline example processor:

```
ADD R1, R2, R3
SUB R4, R1, R5
AND R6, R1, R7
OR R8, R1, R9
XOR R10, R1, R11
```

Data Hazards (2/4)



- The use of results from ADD instruction causes hazard since the register is not written until after those instructions read it.

Data Hazards (3/4)

• Explanation

- **ADD instruction** writes the result in register R1 only at the write back stage, but SUB instruction reads the value during its instruction decode stage. This is called a data hazard.
- **SUB instruction** will read the wrong value and will use it.
- **AND instruction** is also affected by this hazard. The AND instruction that reads the registers in clock cycle 4 will receive the wrong results.

Data Hazards (4/4)

● Explanation

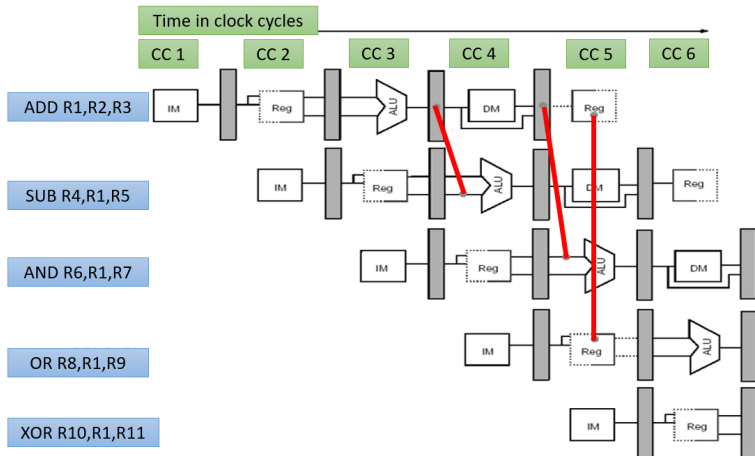
- **OR instruction** can also work fine without incurring a hazard, using a simple implementation technique.
- The technique is to perform the register file reads in the second half of the clock cycle and the writes in the first half.
- **XOR instruction** operates correctly, it reads its inputs (in clock cycle 6) after the ADD has written its result (in clock cycle 5).

Data Hazards – Solution

- The data hazards can be solved by the following three solutions depending on the situation.
 - Forwarding
 - Stalling
 - Compiler Scheduling

Data Hazards – Forwarding

- Eliminate the stalls for the hazard involving SUB and AND instructions using a technique called forwarding



Data Hazards Classification (1/4)

- Depending on the order of read and write access in the instructions, data hazards could be classified as three types.
- Consider two instructions i and j , with i occurring before j .
- Possible data hazards are:
 - **RAW (Read After Write)**
 - **WAW (Write After Write)**
 - **WAR (Write After Read)**

Data Hazards Classification (2/4)

- **RAW (Read After Write)** j tries to read a source before i writes to it, so j incorrectly gets the old value.

Data Hazards Classification (3/4)

- **WAW (Write After Write)** j tries to write an operand before it is written by i.
- The writes end up being performed in the wrong order, leaving the value written by i rather than the value written by j in the destination.

Data Hazards Classification (4/4)

- **WAR (Write After Read)** j tries to write a destination before it is read by i, so the instruction i incorrectly gets the new value.

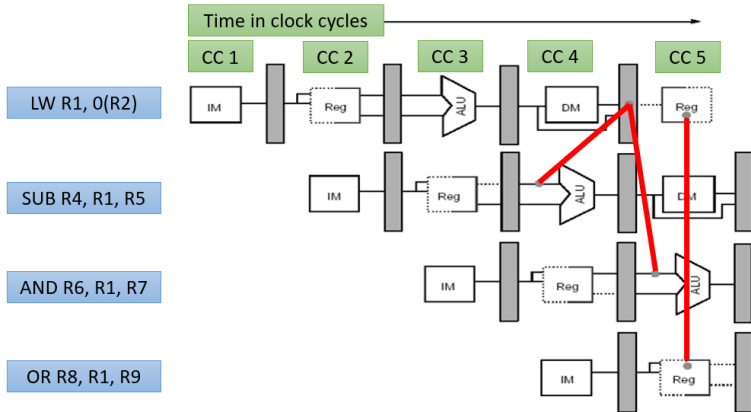
Data Hazards – Requiring Stalls (1/3)

- Unfortunately not all data hazards can be handled by forwarding.
- Consider the following sequence:

```
LW R1, 0(R2)
SUB R4, R1, R5
AND R6, R1, R7
OR R8, R1, R9
```

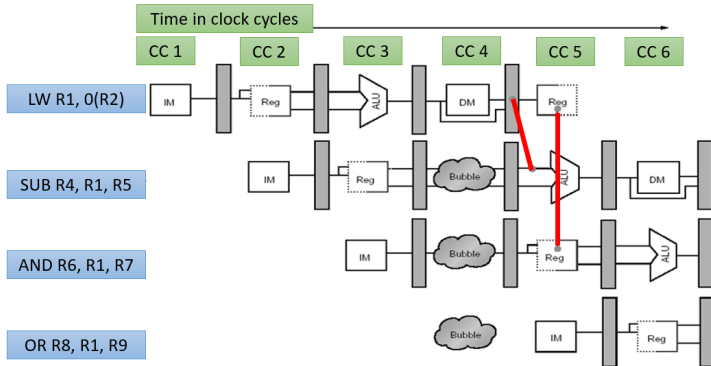
- The problem with this sequence is that the Load operation will not have data until the end of memory stage.

Data Hazards – Requiring Stalls (2/3)



- The load instruction can forward the results to AND and OR instruction, but not to the SUB instruction since that would mean forwarding results in “negative” time .

Data Hazards – Requiring Stalls (3/3)



- The load interlock causes a stall to be inserted at clock cycle 4, delaying the SUB instruction and those that follow by one cycle.
- This delay allows the value to be successfully forwarded onto the next clock cycle.

Data Hazards – Compiler Scheduling


- Rather than just allow the pipeline to stall, the compiler could try to schedule the pipeline to avoid the stalls, by **rearranging the code**.
 - The compiler could try to avoid generating the code with a load followed by an immediate use of the load destination register.
 - This technique is called **pipeline scheduling** or **instruction scheduling** and it is a very used technique in modern compilers.

Control Hazards (1/2)

- **Control Hazards:** It attempts to make branching decisions before branch condition is evaluated.
- Most of the work for a branch computation is done in the execution stage.
- The branch target address is computed.
- The source registers are compared by the ALU, and the Zero flag is set or cleared accordingly.

Control Hazards (2/2)

- The branch decision cannot be made until the end of the execution stage.
- But we need to know which instruction to fetch next, in order to keep the pipeline running!
- This leads to what is called a control hazard.




Control Hazards – Solutions (1/6)

- The control hazards are handled by the following possible ways.
 - Stalls
 - Branch Prediction

Control Hazards – Solution (2/6)

Stalls

- Stalling is one possible solution.



- Here, we just stall until cycle 4, after we do make the branch decision.

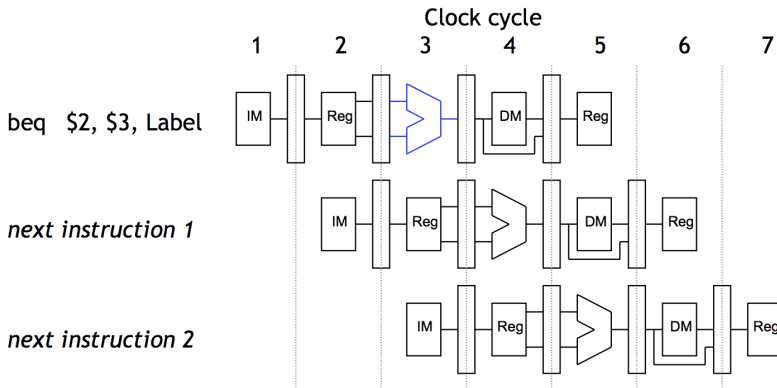
Branch Prediction

- Another approach is to guess whether or not the branch is taken.
 - In terms of hardware, it's easier to assume the branch is not taken.
 - This way we just increment the PC and continue execution, as for normal instructions.
 - If we are correct, then there is no problem and the pipeline keeps going at full speed.

Control Hazards – Solution (4/6)

Example:

- We can get the branch equal results in cycle 3 and possible to make branch decision in cycle 4.

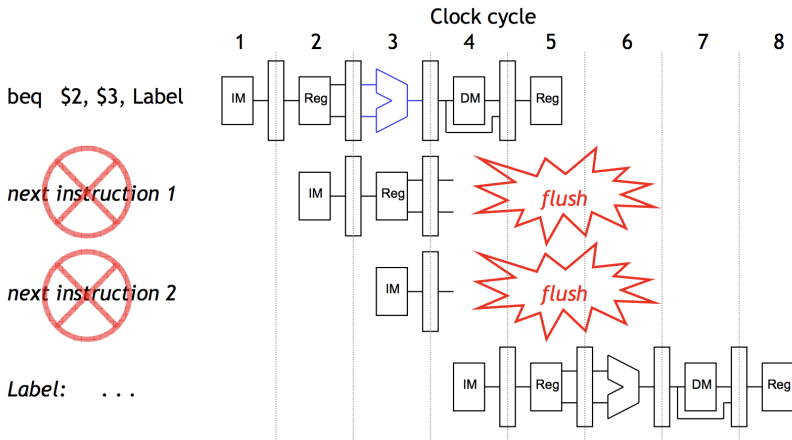


Control Hazards – Solution (5/6)

- If our guess is wrong, then we would have already started executing two instructions incorrectly.
- We need to discard, or flush, those instructions and begin executing the right ones from the branch target address, **Label**.
- It is also known as **branch misprediction**.

Control Hazards – Solution (6/6)

Branch Misprediction



Branch Prediction (1/3)

- Ideal pipelined processor: $CPI = 1$
- Branch misprediction increases CPI
- Overall, branch prediction is worth it.
 - Mispredicting a branch means that two clock cycles are wasted.
 - But if our predictions are even just occasionally correct, then this is preferable to stalling and wasting two cycles for every branch.
 - It is known as **Static Branch Prediction**

Branch Prediction (2/3)

- All modern CPUs use branch prediction.
 - Accurate predictions are important for optimal performance.
 - Most CPUs predict branches dynamically – statistics are kept at runtime to determine the likelihood of a branch being taken.
 - It is known as **Dynamic Branch Prediction**

Branch Prediction (3/3)

- The pipeline structure has a big impact on branch prediction.
 - A longer pipeline may require more instructions to be flushed for a misprediction, resulting in more wasted time and lower performance.
 - We must be careful that instructions **do not modify registers or memory before they get flushed.**

Smart Solution


- We can actually decide the branch a little earlier, in instruction decoding stage instead of execution stage.
- We can add a small comparison circuit to the instruction decode stage, after the source registers are read.
- Then we would only need to flush one instruction on a misprediction.

Optimize Code

- When working with conditional code (if-else statements), we often know which branch is true and which is not. If compiler knows this information in advance, it can generate **most optimized code**.
- For Example: **likely** and **unlikely** macros help the compiler know whether an if is usually going to be entered or skipped.
- Using it results in some **performance improvements**.

Control Hazards – Alternative Solution (2/2)

Smart Solution




- We must flush one instruction, if the previous instruction is BEQ and its two source registers are equal.
- Flushing introduces a bubble into the pipeline, which represents the one cycle delay in taking the branch.

A Superscalar CPU

- Duplicating hardware in one pipe stage would not help
 - For Example: have 2 ALUs
 - The bottleneck moves to other stages
- **Is superscalar good enough?**
 - A superscalar processor can fetch, decode, execute and write back 2 instructions in parallel
 - Can execute only independent instructions in parallel
 - But ... adjacent instructions are usually dependent
- **Solution: out-of-order execution**
 - Execute instructions based on “data flow” rather than program order
 - Still need to keep the semantics of the original program

Out-of-order execution – General Scheme



- It creates new dependencies WAR and WAW.
- These are false dependencies
- There is no missing data
- Still prevent executing instructions out-of-order
- Solution:** Register Renaming

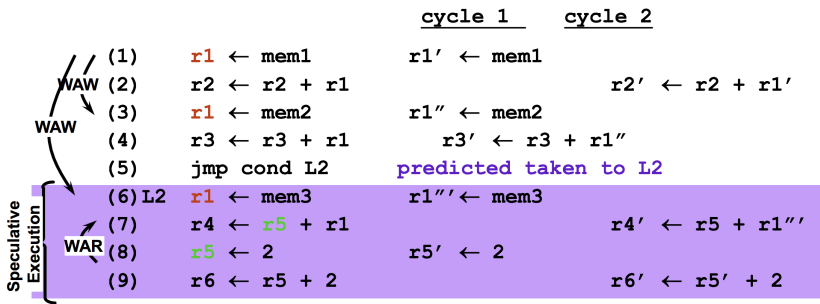
Executing Beyond Branches

- So far we do not look for instructions ready to execute beyond a branch
 - Limited to the parallelism within a basic-block
- We would like to look beyond branches
 - But what if we execute an instruction beyond a branch and then it turns out that we predicted the wrong path?
- **Solution: Speculative Execution**

Speculative Execution

- Hold a pool of all not yet executed instructions
- Fetch instructions into the pool from a predicted path
- Instructions for which all operands are ready can be executed
- An instruction may change the processor state (commit) only when it is safe

Speculative Execution – Example



- Instructions 6-9 are speculatively executed
 - If the prediction turns wrong, they will be flushed
- If the branch was predicted taken
 - The instructions from the other path would have been speculatively executed

Summary

- Hazards
- Types of Hazards
- Structural Hazards
- Structural Hazards Solutions
- Data Hazards
- Data Hazards Solutions
- Control Hazards
- Control Hazards Solutions
- Superscalar CPU

Acknowledgements

- This lecture was created and maintained by Muhammad Fahim, Giancarlo Succi and Alexander Tormasov